

Detecting marine heatwaves with sub-optimal data

Robert W. Schlegel(1,2), Eric C. J. Oliver(1), Alistair J. Hobday(3), Albertus J. Smit(2)

April 30th, 2018

1: Department of Oceanography, Dalhousie University, Halifax, Nova Scotia, Canada 2: Department of Biodiversity and Conservation Biology, University of the Western Cape, Bellville, South Africa 3: CSIRO Oceans and Atmosphere, Hobart, Tasmania, 7000, Australia

Abstract

Marine heatwaves (MHWs), or prolonged periods of anomalously warm sea water temperature, have been increasing in duration and intensity globally for decades. However, there are many coastal, oceanic, and polar regions where our ability to detect MHWs is uncertain due to the unavailability of high-quality data. Here we investigate the effect that short time series length, missing data, or linear decadal temperature trends may have on the detection of MHWs. We show that MHWs detected in time series as short as 10 years did not have durations or intensities significantly different from events detected in the standard length 30 year time series, but the accurate identification of temperature thresholds could be impaired when fewer than 15 years of data were used. We also show that the output of the MHW algorithm for time series missing less than 20% data did not differ appreciably, and that this could be increased to 40% when gaps were filled with linear interpolation. Linear decadal trends of $0.05 - 0.15^{\circ}\text{C}/\text{dec}$ could lead to inaccurate creation of seasonal climatologies, but this did not impact accurate MHW detection. The percentage of missing data in a time series was determined to have the largest effect on the detection of MHWs, but was also the easiest to correct for. Time series length had less of an effect on MHW detection, but was more difficult to correct for. We provide suggestions for best practices to improve the accuracy of MHW detection with sub-optimal time series on a global scale and specific case studies of three notable MHWs from the literature.

Introduction

The idea of locally warm seawater being problematic is not a novel concept. We have known for decades that seemingly transient warm water occurrences in the ocean could result in major impacts (e.g. Baumgartner 1992, @Salinger2016). The study of the effects of anomalously warm seawater temperatures began in earnest in the early 1980's when research into the ENSO phenomenon began (e.g. Philander 1983). After the 1980's, researchers began noticing that warm water events were becoming more frequent and problematic, but it wasn't until 2018 that this was demonstrated with global observations (Oliver et al. 2018).

In order to quantify the increased occurrence and severity of these events it was necessary to develop a methodology that would be inter-comparable for the entire planet. This was accomplished in 2016 after the International Working Group on Marine Heatwaves (marineheatwaves.org) initiated a series of workshops to address this issue (Hobday et al. 2016). This definition for anomalously warm seawater events, known as marine heatwaves (MHWs), has seen wide-spread and rapid adoption due to ease of use and applicability to any part of the globe. One problem with this algorithm that has not yet been addressed is the assumption that a researcher has access to the highest quality data available when detecting MHWs. In the context of MHW detection 'high quality' is a daily time series with no missing data that is at least 30 years in length. To avoid contention on the use of the word 'quality', time series that meet the aforementioned standard are referred to here as 'optimal', whereas those that do not meet some part of the standard are referred to as 'sub-optimal'.

Most remotely-sensed data, and more recently output from ocean models and reanalyses, consist of over 30 years of data and utilise statistical techniques to fill gaps in the time series from a number of environmental and technical sources. This means that these data are considered optimal for MHW detection however, they still have issues (e.g. land bleed and incorrect data flagging) and so it may be recommended that researchers utilise sub-optimal data when possible, such as sporadically collected *in situ* time series. Coastal areas are often poorly sampled yet are the most susceptible to the impacts of MHWs (e.g. Smale et al. 2019) so it is necessary to address the issues that using these data may have on the detection of MHWs.

This paper seeks to understand the limitation of using sub-optimal data for the detection of MHWs. Of primary interest are three key challenges: 1) The use of short time series, 2) the use of time series with missing data, 3) the use of time series in areas with large long-term temperature trends. We will use a combination of reference time series from specific locations and global data to address these issues. The effects of the three sub-optimal data challenges on the detection of MHWs are quantified in order to provide researchers with the level of confidence they may express in their results. Where possible, best practices for the correction of these issues are detailed.

Defining marine heatwaves

The definition used in this paper for a MHW is “a prolonged discrete anomalously warm water event that can be described by its duration, intensity, rate of evolution, and spatial extent.” (Hobday et al. 2016). This qualitative definition is quantified with an algorithm that calculates a suite of metrics. These metrics may then be used to characterise MHWs and to effectively compare them against known ecological/financial impacts. The calculation of these metrics is made possible by first determining the mean and 90th percentile temperature for each of the 366 calendar day-of-year (doy) in a time series. The mean doy temperatures, which also represent the seasonal signal in the time series, provide the expected baseline temperature whose daily exceedance is used to calculate the intensity of MHWs. The 90th percentile doy temperatures serve as the threshold that must be exceeded for 5 or more consecutive days for the anomalously warm temperatures to be classified as a MHW and for the calculation of the MHW metrics to begin.

In this paper we will focus on the two metrics that most succinctly summarise a MHW. The first metric, *duration (days)*, is defined as the number of days that the temperature remains at or above the 90th percentile threshold without dipping below it for more than 2 consecutive days. The duration of an event is the best single measurement of the chronic stress that a MHW may inflict upon a target species or ecosystem. The second metric, *maximum intensity (°C)*, is the single warmest day during the event and is calculated by subtracting the mean doy temperature on that day from the recorded temperature. This metric is the best single representation of acute stress. There are many other MHW metrics and the full explanation for them may be found in Table 2 of Hobday et al. (2016).

Hobday et al. (2018) extended the MHW definition to include a categorisation scheme based on the intensity of an event. These categories were: I Moderate, II Strong, III Severe, and IV Extreme. The category of an event is determined by how many times the maximum intensity of the MHW is a multiple of the difference between the mean and 90th percentile doy temperatures (Figure 1). For example, if the difference between the mean and 90th percentile doy temperatures on the warmest day of a MHW 1.5°C, and the temperature recorded on that warmest day was 3°C warmer than the mean doy temperature for that day, this would be considered a category II (Strong) MHW. Were the maximum temperature recorded at 4.5°C, this would then be classified as a category III Severe MHW. To provide a more robust qualification of a MHW, the categories are also calculated for each day of a MHW to provide a proportion of the days during which the event was within each of the categories.

An additional advantage in the use of the Hobday et al. (2016) and Hobday et al. (2018) approach is that it has been developed for python (<https://github.com/ecjoliver/marineHeatWaves>), R (Schlegel and Smit 2018), and MATLAB (Zhao and Marin 2019). For this analysis we compared the R and python default outputs, assessed how changing the arguments affected the results, and compared the other functionality provided between the two languages. While some style differences exist between the added functionality of

the languages, the core climatology outputs are identical to within < 0.001 °C per day-of-year (doy). An independent analysis of the Python and MATLAB results also confirmed that they were functionally identical (pers. com. Zijie Zhao; MATLAB distribution author).

What are optimal data for detecting marine heatwaves?

Hobday et al. (2016) stated that optimal data for detecting MHWs have the following properties: 1) the time series must be at least 30 years in length, 2) be quality controlled, 3) be of the highest resolution possible, and 4) *in situ* data should be used to compliment remotely sensed data where possible. Whereas the authors did not specifically state that time series must not contain large proportions of missing data, it can be inferred from the aforementioned requirements. There are a number of methods within the already existing tools for detecting MHWs that can address these concerns and we will lay them out here. An issue not discussed in Hobday et al. (2016) is the effect of long-term trends on the accurate detection of events. Oliver et al. (2018) have shown how dominant the climate change signal can be in the detection of events and we seek to quantify this effect here.

A time series with a sub-optimal length may impact the detection of MHWs by negatively affecting the creation of the daily climatology relative to which MHWs are detected in two primary ways. The first is that with fewer years of data to draw from, the presence of an anomalously warm or cold year will have a larger effect on the climatology than with a sample size of 30 years. The second cause is that because the world is generally warming (Pachauri et al. 2014), the use of a shorter time series will almost certainly warm bias the results.

The climatology derived from a time series serves two main roles (Organization 2017); 1) it serves as a ‘benchmark’ relative to which past and future measurements can be compared, and against which anomalies can be calculated, 2) it reflects the typical conditions likely to be experienced at a particular place at a particular time. The WMO Guide to Climatological Practices (Organization 2011) stipulate that daily climatologies (which they call ‘climate normal’) must be based on the most recent 30-year period that ends on a complete decade (currently 1981 – 2010). It is from this WMO guideline that the optimal length for MHW detection was derived.

Some remotely sensed products suffer from ‘gappiness’ that result in missing data being introduced. This may be due to cloud cover, the presence of sea ice, unsuitable sea states, etc., which become more prevalent at smaller scales, particularly nearer the coast. Some products smooth out these influences, but this results in smoothed SST fields that mask some of the small-scale spatial variation in surface temperatures. Other times they rely on blending with data from other products, which may have its own suite of consequences. This is why the use of imperfect *in situ* collected time series may still be encouraged in certain situations. These data are however also prone to large gaps and so the problems these data face with regards to accurate event detection are generally uncertain.

Methods

To quantify the effects that time series length, missing data, and long-term trends have on MHW detection we will focus on the following three outputs created by the MHW detection algorithm:

- 1) The climatologies derived from the daily SST records, which include both the seasonally-varying mean and 90th percentile threshold.
 - These are not a part of the MHWs themselves, but are necessary for their detection.
- 2) The MHW event itself, which is defined by the metrics given in Table 2 of Hobday et al. (2016).
 - We chose here to focus on only the duration (days) and maximum intensity (°C) metrics in order to keep the results manageable.
- 3) The proportion of days of the event that are within the different categories.
 - These are a more qualitative result that may be more applicable to a broader audience.

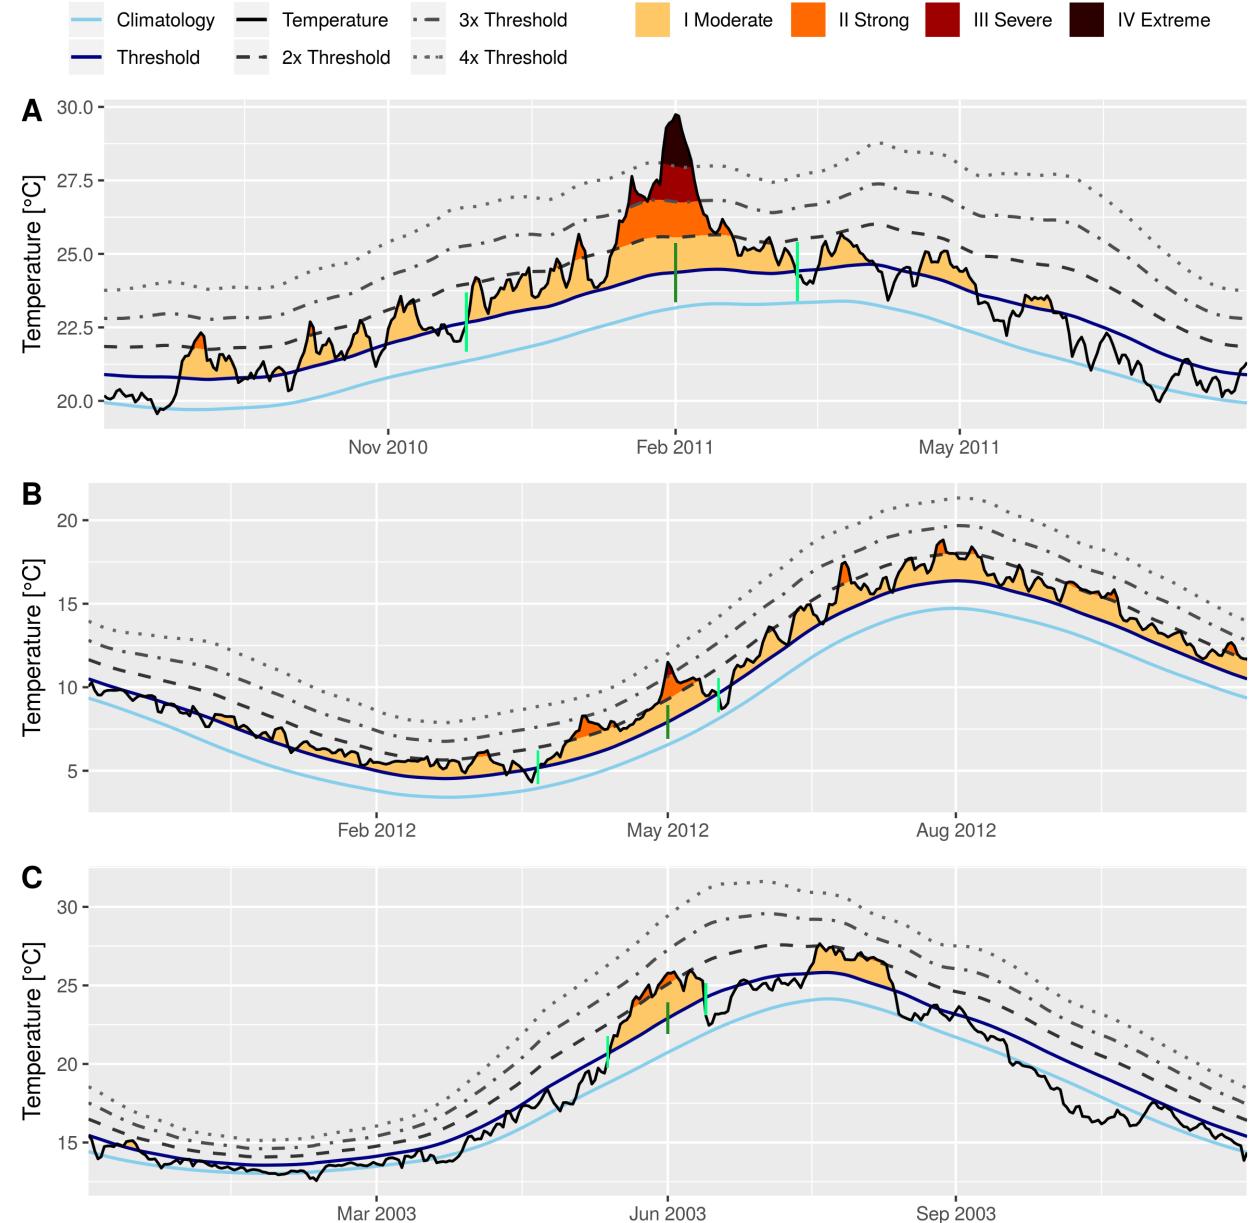


Figure 1: Figure 1: The focus marine heatwaves (MHWs) from the three reference time series A) Western Australia, B) Northwest Atlantic, and C) Mediterranean. The various thresholds for the different MHW categories are shown as grey dotted lines and the proportions of the MHW within each category are filled with the corresponding colours seen in the legend at the top of the figure. Each panel is centred around the peak date of the focus MHW. The peak date is further highlighted by a dark green vertical segment. The beginning and end of each event are demarcated with light green vertical segments.

With these three focal items defined, we will then use the following three questions to frame the results:

- 1) How sub-optimal can data be before any of the above three items become significantly different from those calculated with an optimal time series?
 - For example, how short may a time series be before the climatology becomes significantly different from the same climatology derived from the full 30 year time series?
- 2) What amounts of uncertainty are introduced into the results from the increasingly sub-optimal data?
 - For example, when 20% of data are missing, what should a user expect the standard error around the duration of a MHW to be compared to that same MHW when detected in a time series missing no data?
- 3) Are the error rates introduced by sub-optimal data for the event metrics the same/similar everywhere in the world, or do they differ based on some observable pattern/known oceanographic feature(s)?
 - For example, when the length of a time series is shortened to 10 years in an eastern boundary upwelling system (EBUS), does the effect this have on the maximum intensity of the events differ from the same shortening on a time series in a western boundary current (WBC)?

To answer these three questions we will use the remotely sensed NOAA OISST dataset (Reynolds et al. 2007, Banzon2016). This daily remotely-sensed global SST product has a 1/4 degree spatial resolution. The first complete year of data available is 1982, meaning that we must deviate slightly from the WMO standard for daily climatology creation by setting our reference period at 1982 – 2011 unless otherwise noted.

The first two questions posed above will be answered using the three reference time series from Hobday et al. (2016). These time series are taken from the coast of Western Australia (WA; Figure 1A), the northwest Atlantic Ocean (NWA; Figure 1B), and the Mediterranean Sea (Med; Figure 1C). These time series are used here for ease of reproducibility and because they each contain a MHW that has been the focus of multiple publications.

For the third question posed above we will use the entire global NOAA OISST product. Each pixel in this dataset will have the single largest event in the most recent ten years of data (2009 – 2018) identified and as the different sub-optimal tests are performed the effect this has on the event metrics will be recorded so that the relationship they have with sub-optimal data may be quantified. For this test we will not be removing the long-term trend in the data as we want to see what the real-world pattern in the data are.

The outputs of the MHW detection algorithm from the reference time series with different sub-optimal challenges will be compared against the same optimal reference time series (i.e. 30 year length, no missing data) with a Kolmogorov-Smirnov (KS) test. This test looks for differences in the continuous distribution of values between two sets of data, rather than testing for differences of central tendency (e.g. *t*-test or ANOVA), and provides a *p*-value that indicates the probability that the two distributions being compared have been drawn from the same pool of data. It was decided not to test for central tendency for two reasons; the first being that the assumption of normality for the values in the outputs was usually violated, and the second was that we do not want to know how sub-optimal data affect the central tendency of the results, but rather how they affect the distribution of the results. For example, does a 15 year time series produce a larger number of short events than a 30 year time series? To this end we are also not interested in rejecting a null hypothesis that the outputs from the sub-optimal data are the same as the optimal data based on a *p*-value of 0.05 or less. This is in part because testing for null hypotheses in this way is becoming increasingly discouraged (Wasserstein, Schirm, and Lazar 2019), but also because we want to show what the probability is that results from a time sub-optimal time series may be different from an optimal time series. That being said, we will still highlight comparisons that generate a *p*-value of 0.05 or less.

Because it would not be a robust test of the effects of sub-optimal data on MHW detection to use only three time series in this way, the order of the years within each of the three reference time series were randomly re-sampled and recombined 100 times, ensuring that one of the re-samples maintained the original order of the reference time series. This 100 fold increase to the available dataset will allow for a better estimate of the error that sub-optimal data introduce into MHW detection. We chose this method instead of creating artificial time series with comparable auto-correlation structures as it ensured that the large historical MHWs present in the reference time series would still be used in the calculations as these are an important reason why these time series were chosen.

Lastly, because the effects of time series length and long-term trends are to be quantified separately, it is necessary to de-trend the time series before beginning to control for the sub-optimal challenges. This de-trending is performed by fitting a simple linear model to each of the re-sampled time series and then removing it from the data. The following three sub-sections describe how the three sub-optimal time series challenges will be controlled for.

Controlling for time series length

There are currently 37 complete years of data available in the NOAA OISST dataset (1982 – 2018). In order to determine the effect that time series length has on the three MHW detection algorithm outputs, we will systematically shorten each of the 100 re-sampled reference time series, one year at a time from 37 years down to 10 years (2009 – 2018), before running the MHW detection algorithm. The three different outputs (climatologies, event metrics, and categories) for each one year step for each of the re-sampled time series will then be compared against the output from the optimal 30 year version of that same time series using a KS test.

In order to ensure equitable sample sizes we will only be comparing the MHW metrics and categories for events detected within the last 10 years of each test as this is the period of time during which all of the different tests overlap. This is also why we have limited the shortening of the time series lengths to 10 years, so that we could still have a reasonable sample size to draw from.

Because the lengths of the time series were being varied, and were usually less than 30 years in length, it was also necessary that the climatology periods vary likewise. To maintain as much consistency as possible across the results we used the full range of years within each shortened time series to determine the climatology. For example, if the time series had been shortened from 37 to 32 years (1987 – 2018), the 32 year period was used to create the climatology. If the shortened time series was 15 years long (2004 – 2018), this base period was used. The control time series were those with a 30 year length ending in the most recent full year of data available (1989 – 2018). Note that due to necessity this differs from the climatology period of 1982 – 2011 used for the other tests outlined below.

The *a-priori* fix proposed to address the issue of short time series length is to use a different climatology estimation technique. The option currently available within the MHW detection algorithm is to expand the window half width used when smoothing the climatology. Other techniques, such as harmonic regression/Fourier analysis, would have a similar effect but are not used here in favour of the methodology available within the MHW algorithm.

Controlling for missing data

In order to determine how much random missing data effect the outputs of the MHW algorithm, we will randomly removed 0 – 50% of the data in 1% steps from each of the re-sampled time series before running the MHW algorithm on each step. The optimal time series against which the various outputs are compared via a KS test will be the same re-sampled time series with 0% missing data.

The *a-priori* fix for the issue of missing data in the time series is to linearly interpolate over any gaps. There are many methods of interpolation (imputing) gaps in time series, but we choose linear interpolation because of its simplicity and because it is already available in the software implementations of the MHW algorithm.

Controlling for long-term trends

To quantify the effect of long-term (secular) trends on the outputs of the MHW algorithm we added linear decadal trends of 0.00 – 0.50°C/dec in 0.01°C steps to each of the re-sampled time series. The difference this caused in the outputs was quantified with the same tests as for length and missing data. The optimal time series used a control for the KS comparisons were those with no added trend.

There is no proposed *a-priori* method to correct for the added linear decadal trend in these data as this would be to simply not add it. Rather it is proposed that the relationship between the slope of the added trend and the results it has on the outputs of MHW algorithm be documented to determine if a predictable relationship may be used to correct the results *post-hoc*.

Results

Time series length

Shortening the lengths of the re-sampled time series had a noticeably negative effect on the comparability of the outputs of the MHW algorithm (Figure 2). We see that the climatology outputs were affected the most, and the category proportions affected the least. Changing the length of a time series lowered the mean probability (*p*-value) from the 100 re-sampled tests for each reference time series, but was accompanied by a high level of variance.

With the exception of the Western Australia (WA) time series we see that there is no point at which any of the outputs from the MHW algorithm on shortened time series became significantly different from the 30 year control time series. The WA time series, which is characterised by its large inter-annual variability, only shows significantly different threshold climatologies on average when 14 years of data or fewer are used (Figure 2A). The seasonal climatology does not differ significantly on average until 11 years of data or fewer are used.

It is important to note that increasing the climatology period longer than 30 years has almost as rapid an effect on creating dissimilar outputs as using fewer years of data does. This was an unexpected result that stresses the importance of adhering to the WMO standard as closely as possible to ensure the comparability of results.

When we look at the effect that shortening time series length has on the duration and maximum intensity of the focus MHWs from the original data (not re-sampled) from each reference time series we see that both values tend to decrease (Figure 3). This is because the shortening of a time series tends to increase the mean climatology, so the shorter a time series becomes, the lower the maximum intensity and shorter the duration of the MHWs become relative to the increasing mean climatology (Figure 3; left column bottom and middle panels). We also see that the Western Australia (WA) MHW is always being divided into at least two or more separate events due to the rising mean climatology (Figure 3; left top panel). The Mediterranean (Med) MHW isn't affected much by changes in time series length as this MHW has a much more even rise and fall above the 90th percentile threshold (Figure 1C).

The gently increasing maximum intensities caused by increasing the lengths of the time series were also generally seen to occur throughout the oceans (Figure 4). The median rate of change caused by increasing time series length from 10 to 30 years is seen to be 0.5% per year. This means that, on average, a MHW detected in a 10 year time series will have a max. intensity about 10% cooler than a MHW detected in a 30 year time series. This is a very small margin and shows the robustness of the MHW detection methodology. We also see that the only areas that show MHWs decreasing in maximum intensity are most of the Southern Ocean and some parts of the open ocean in the Pacific and Atlantic.

Remember that the long-term trend in these data were not removed beforehand so it is not surprising that increasing the length of the time series into the past (where the data are cooler on average) will reduce the mean climatology and therefore increase the maximum intensity of the detected event. The relationship between warming or cooling maximum intensity and decadal trends in temperature (i.e. climate change) is significant (*p*-value < 0.001, $R^2 = 0.33$). We also see that areas with perennial ice coverage, and western boundary currents (WBCs), tend to show greater rates of change. This is likely due to these areas having larger amounts of variance, in addition to a stronger decadal warming trend.

Looking at the effect of time series length on the duration of MHWs around the globe (Figure 5) we see a similar pattern to the effect on maximum intensity (Figure 4). The median increase is 1.4% per year over the

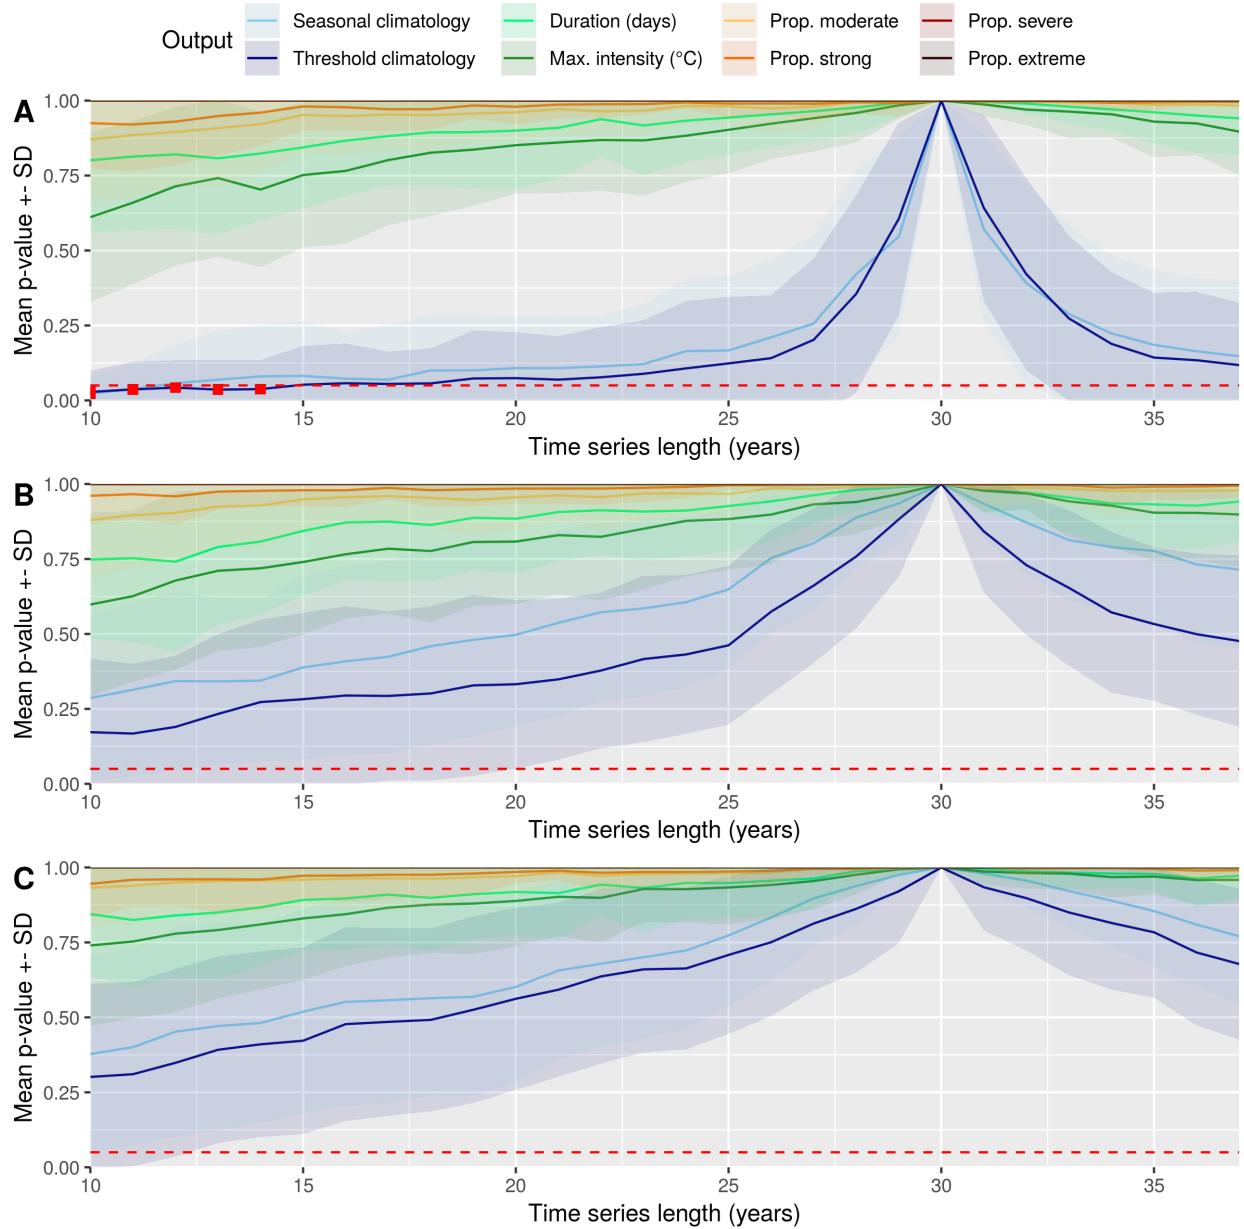


Figure 2: The results from Kolmogorov-Smirnov (KS) tests on the similarity of the outputs from the MHW algorithm with optimal data against sub-optimal time series of differing lengths. The climatology outputs are shown in blue, the event metrics in green, and the category proportions in yellow-red. The solid lines show the mean p -value from the tests on the 100 re-sampled time series for each step. The coloured ribbons show one standard deviation (SD) in the p -values for each step. The results for each reference time series are A) Western Australia (WA), B) Northwest Atlantic (NWA), C) Mediterranean (Med). The x-axis shows the length of the time series being compared against the 30-year (optimal) data. The y-axis shows the range of mean p -values from 1.0 (exact same) to 0.0 (completely different), with a horizontal dashed red line at $p=0.05$ (statistically significantly different). Any mean values at or below the $p=0.05$ line are highlighted with red squares.

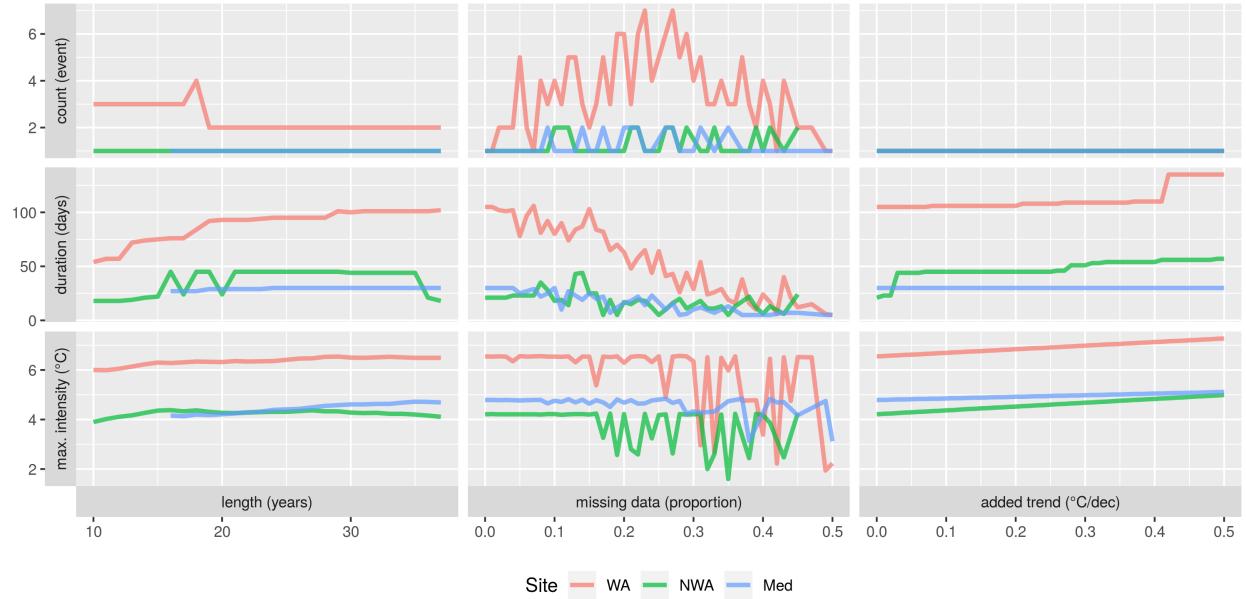


Figure 3: Figure 3: The effect of the three tests (columns) on the MHW metrics and count of events (rows). Each panel has three lines, one for each of the reference time series, shown in the legend at the bottom of the figure. These are the original data, not the randomly re-sampled time series. The lines track the change of just one metric for the focus MHWs seen in Figure 1 as the data are made increasingly sub-optimal, as shown along the x-axes. The y-axes show the unit of measurement for each metric. The top row, “count (event)” shows if the MHW has been divided up into multiple smaller MHWs due to changes in the values along the x-axes.

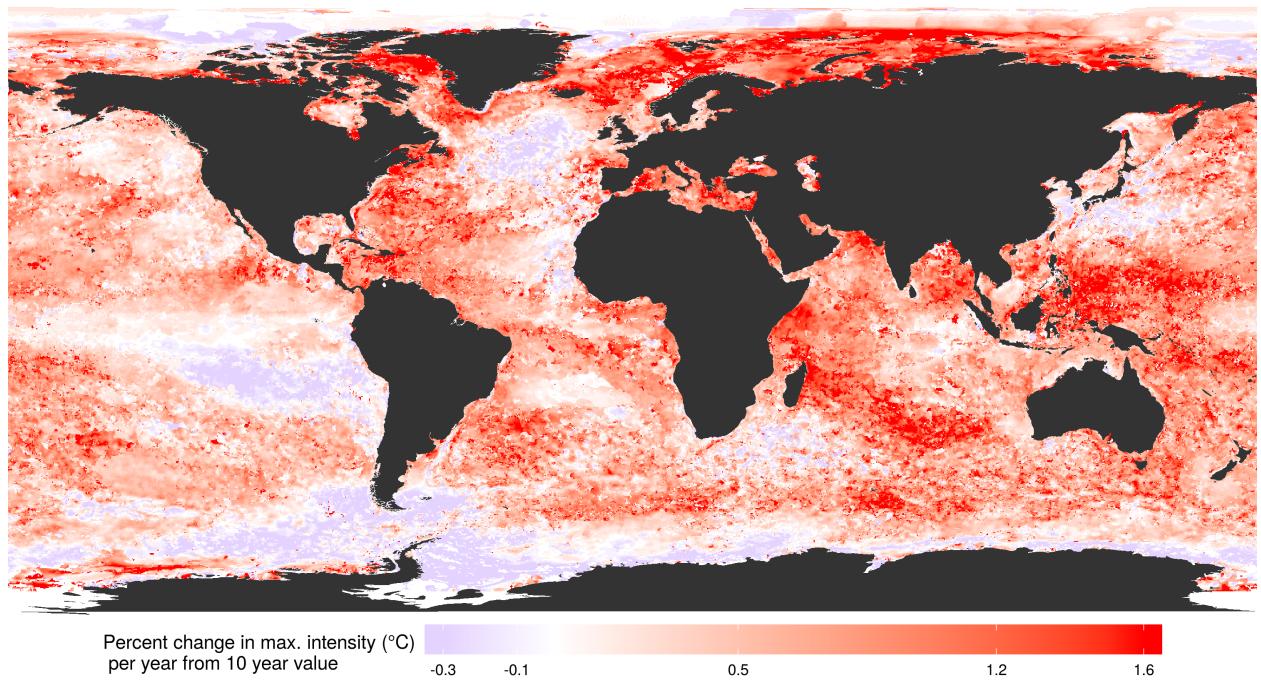


Figure 4: Map showing the percent increase per year in the maximum intensity of the largest MHW detected in the most recent ten years of data when an increasing number of years of data are used for the calculation of the MHW.

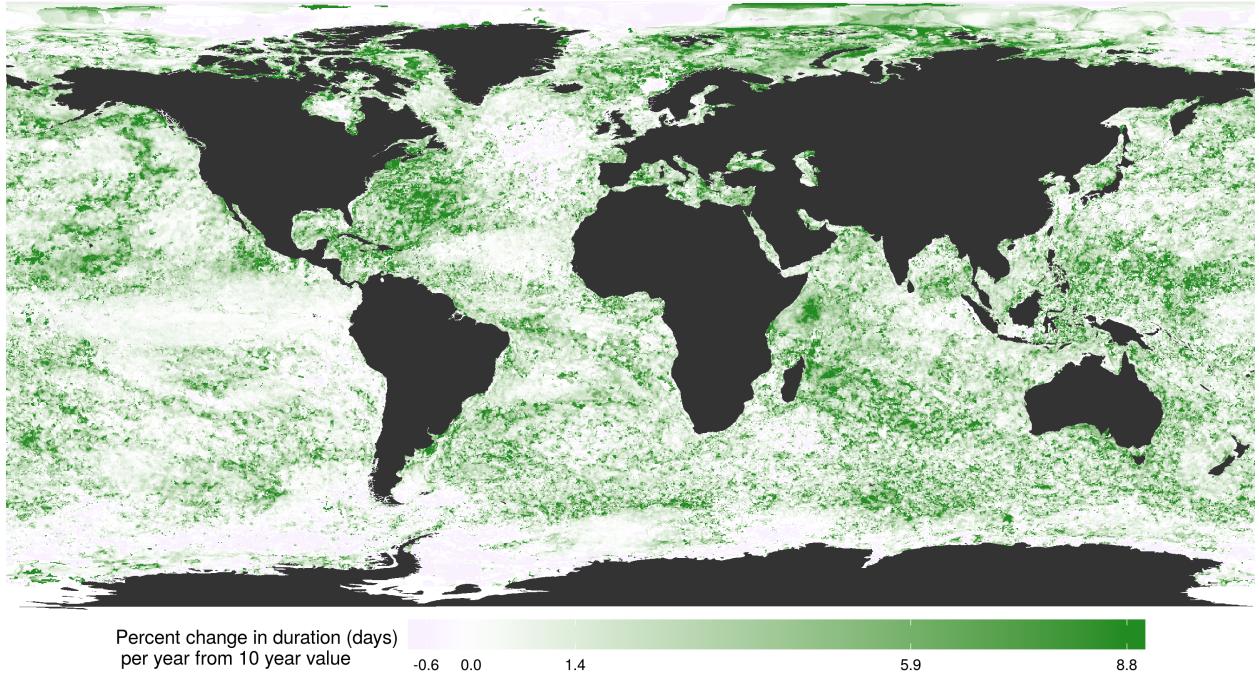


Figure 5: Map showing the percent increase per year in the duration of the largest MHW detected in the most recent ten years of data when progressively more years of data are used for the calculation of the MHW.

duration of the MHW detected with 10 years of data. This is not surprising and supports the observation for maximum intensity.

The fixes proposed for shorter time series may have been beneficial for time series under 15 years in length, but the correction they provided was not consistent. The larger issue cause by a short time series is the amount that the centre of the climatology increases or decreases, more so than the increase in variability caused. This is not something that can be controlled for *a-priori* and is better controlled for in a *post-hoc* manner along the same lines as the proposed fix for decadal trends (see below).

Missing data

The effects of missing data on the outputs of the MHW algorithm are very pronounced. Whereas the changes in time series length may affect the climatologies more rapidly, increases in missing data affect the MHW metrics and the categories much more. The outputs most affected are the threshold climatology, the duration of the MHWs, and the proportions of MHW days in the moderate and strong categories. The maximum intensities of the MHWs are also affected, but at 50% missing data these did not become significantly different from the control time series. The proportion of severe or extreme days were not affected by missing data as they were already so rare or non-existent. The seasonal signal was affected very little by large proportions of missing data.

The effect of random missing data on the single focus MHWs from the three reference time series are very jagged because the missing data at each step was only calculated once. This was done intentionally to highlight the range that this randomness can have on the results as compared to the changes in length (Figure 3; left column). The effect that missing data can have on the MHW metrics depends largely on the shape of the MHW. The WA event has a very pronounced peak (Figure 1A), so when larger proportions of data are

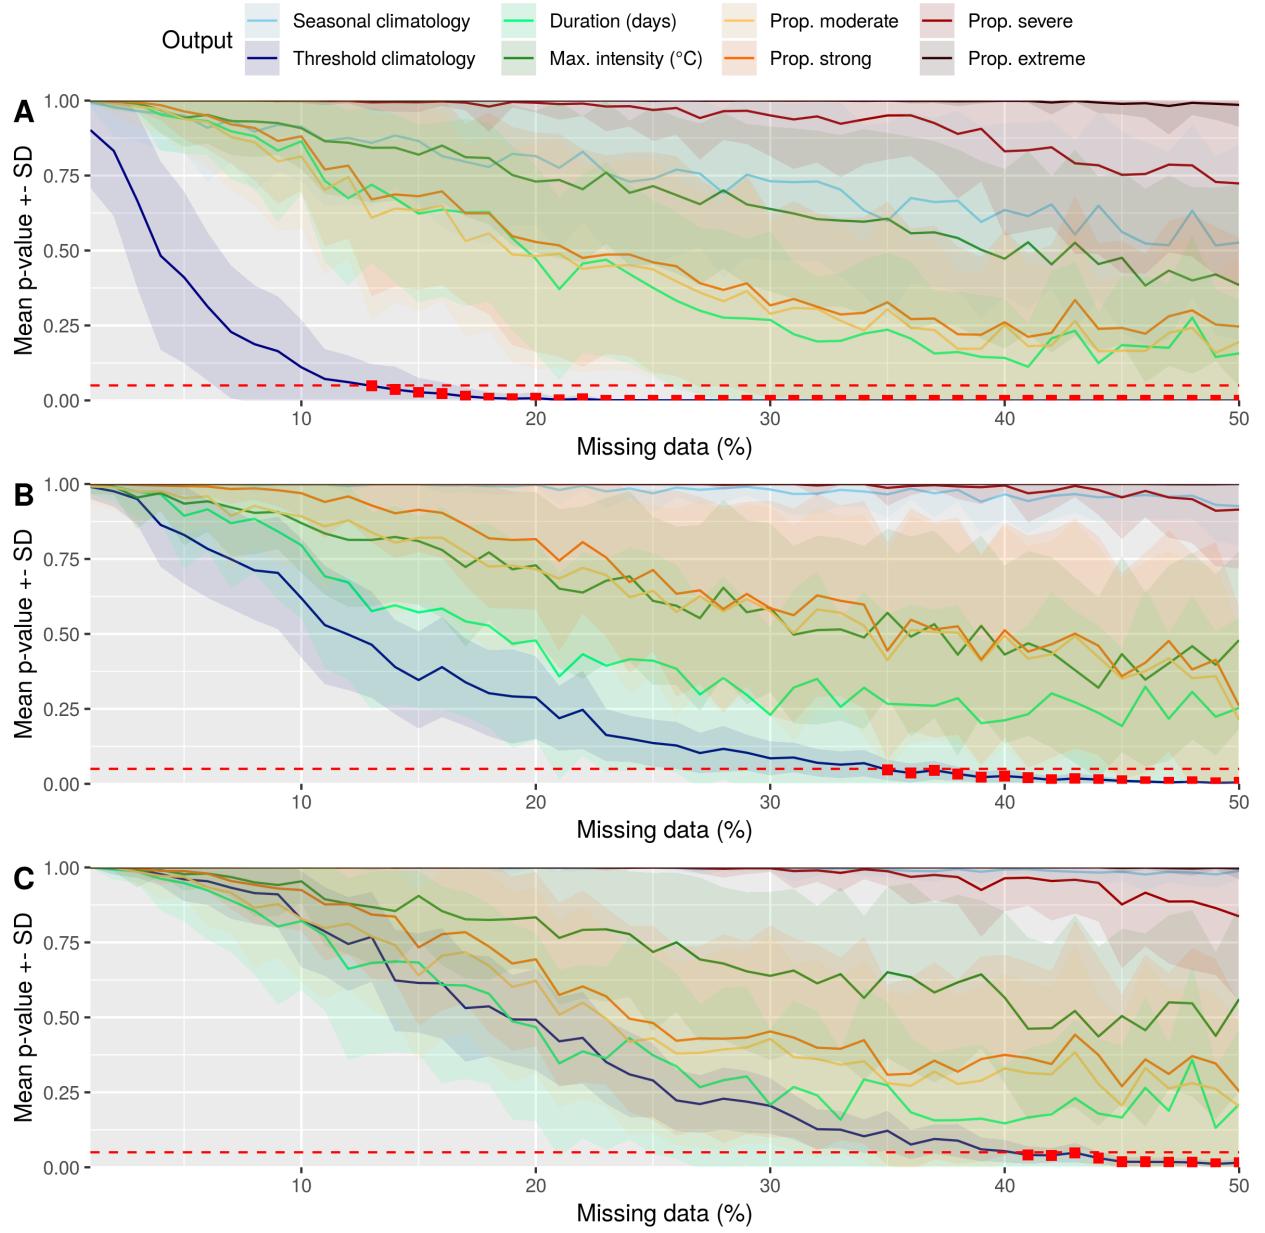


Figure 6: The results from Kolmogorov-Smirnov (KS) tests on the similarity of the outputs from the MHW algorithm with optimal data against sub-optimal time series with increasing percentages of missing data. The climatology outputs are shown in blue, the event metrics in green, and the category proportions in yellow-red. The solid lines show the mean p -value from the tests on the 100 re-sampled time series for each step. The coloured ribbons show one standard deviation (SD) in the p -values for each step. The results for each reference time series are A) Western Australia (WA), B) Northwest Atlantic (NWA), C) Mediterranean (Med). The x-axis shows the percent of missing data in the time series being compared against the complete (0%; optimal) data. The y-axis shows the range of mean p -values from 1.0 (exact same) to 0.0 (completely different), with a horizontal dashed red line at $p=0.05$ (statistically significantly different). Any mean values at or below the $p=0.05$ line are highlighted with red squares.

missing we see how likely it becomes that this peak is not being recorded. The maximum intensity measured in the control time series is 6.5°C , but we see that because very few days of this MHW were so intense, increasing proportions of missing data become more likely to remove these large values. In the NWA event we also see a jagged effect from missing data, though less than the WA event, this is also because of the peak in temperature for this event. The effect on the Med event is the least pronounced. This is because the event does not have one large peak, rather it is more even in its exceedance above the 90th percentile threshold so missing data does not begin to have an appreciable effect on the event until there is an excess of 35% of the data missing.

The duration of the MHWs are all negatively impacted by missing data, with the longer duration MHW (WA) impacted much more than the shorter (NWA and Med) MHWs (Figure 3; centre middle panel). Even though the decrease in duration due to missing data is very rough, we see that it follows a linear trend and can therefore be predicted for within a certain range of error.

For the two shorter MHWs the increase in missing data never divides the event up into more than two separate MHWs (Figure 3; centre top panel). The contiguity of the WA event however is affected greatly by missing data. With just 5% of the data in the time series missing this event was divided into 5 separate events. As missing data increased the count of the divided events tended to also increase up until 27% missing data. At that point the event began to be divided into fewer events again, not because they were forming back together, but rather because there was now too little data to be detecting the splinters being formed off of the main event.

The linear interpolation of missing data was very effective and could potentially allow for the use of time series missing up to 50% of their data (Figure 7), assuming that there is not so much missing data that there are no representative days of the MHW that one may be wanting to study/isolate.

Long-term trends

When adding a linear trend to the re-sampled time series we see that it created statistically significantly different climatologies at an exponential rate (Figure 8). The effect an added decadal trend had on the other outputs of the MHW algorithm was roughly linear, and never produced results significantly different from the control time series. The maximum intensity and duration of events were affected more than the category proportions.

Adding linear long-term trends never caused the focus MHW to be dissected into multiple events (Figure 3; top right panel). The duration of the events are affected differently by the added linear trend. The Med shows practically no effect, the NWA has a very slight increase with a dramatic jump at an added trend of $0.04^{\circ}\text{C}/\text{dec}$, whereas the WA event sees a massive increase due primarily to one large jump at $0.42^{\circ}/\text{dec}$. The effect that the linear trend has on the maximum intensity of each event is a simple linear function of the decadal trend and where in the time series the event occurs. The slope for the increase in maximum intensity for the Med MHW is more shallow than the other two because this MHW occurred in 2003, as opposed to 2010 (WA) and 2012 (NWA).

Discussion

An investigation into the effects that sub-optimal data have on MHW results revealed that there are thresholds within which the outputs of the MHW detection algorithm will remain comparable to results generated by optimal data. Times series longer than about 15 years in length should cause little concern regarding the reliability of the climatologies that are derived from them. The length of a time series has less of an effect on the other outputs of the MHW algorithm, with lengths of 10 years not producing appreciably different outputs in MHW metrics or category proportions. An unexpected result was that increasing the length of a time series longer than 30 years reduced the probability that the outputs would be comparable by as much as shortening the time series did. This means that the common assumptions that using 30 years of data is

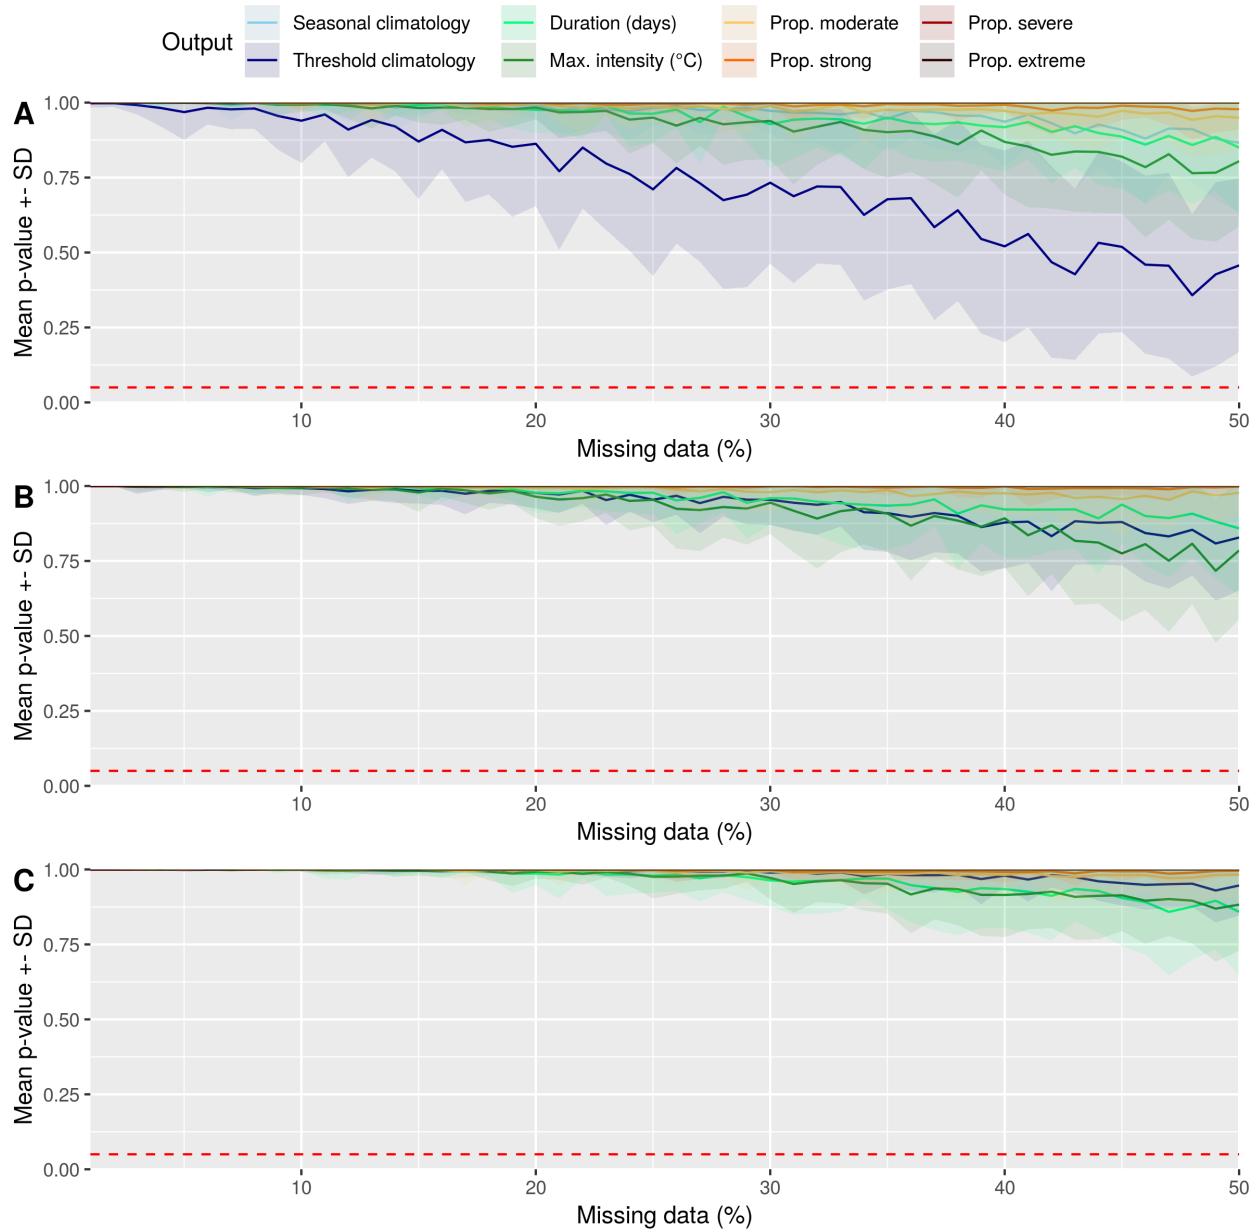


Figure 7: The same information shown in Figure 6, but with the gaps introduced from the random missing data filled via linear interpolation before running the MHW detection algorithm.

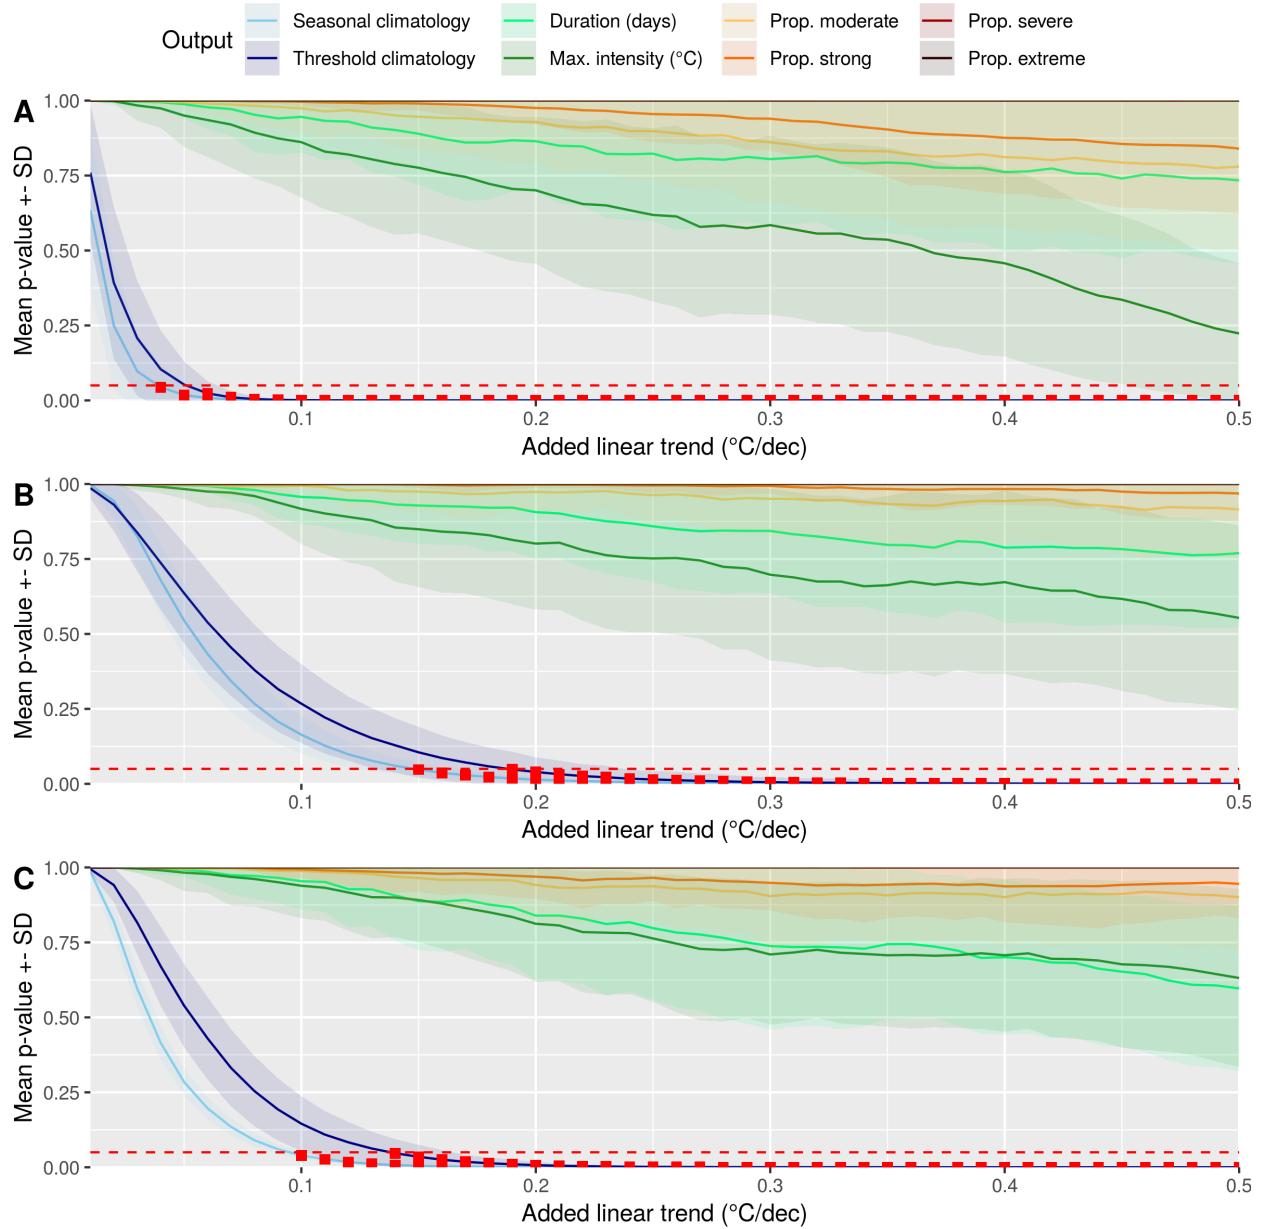


Figure 8: The results from Kolmogorov-Smirnov (KS) tests on the similarity of the outputs from the MHW algorithm with optimal data against sub-optimal time series with increasingly larger linear decadal trends added to them. The climatology outputs are shown in blue, the event metrics in green, and the category proportions in yellow-red. The solid lines show the mean p -value from the tests on the 100 re-sampled time series for each step. The coloured ribbons show one standard deviation (SD) in the p -values for each step. The results for each reference time series are A) Western Australia (WA), B) Northwest Atlantic (NWA), C) Mediterranean (Med). The x-axis shows the decadal trend added to the time series being compared against the flat (0 added trend; optimal) data. The y-axis shows the range of mean p -values from 1.0 (exact same) to 0.0 (completely different), with a horizontal dashed red line at $p=0.05$ (statistically significantly different). Any mean values at or below the $p=0.05$ line are highlighted with red squares.

the same as using > 30 years of data is incorrect. In other words, the 30 year length is often thought of as a minimum length needed to constrain the climatology but we have shown here that using a climatology period greater than 30 years creates different outputs. It is therefore important to stress the adherence to the WMO standards for climatology periods as closely as possible. Increased smoothing of the climatologies derived from shortened time series was not an effective fix to the other outputs of the MHW algorithm. In the global analysis we did see that there is a relationship between decadal trend in seawater temperature and the increase in the duration and maximum intensity of events detected within the most recent decade of data. This can be used to infer a likely correction for the resultant MHW metrics.

The MHW algorithm proved to be resilient to missing data and so long as one does not have particularly large gaps (e.g. greater than a week at a time), time series missing as much as 20% of their data may be used without concern. Greater amounts of missing data could still be used with some caution as the outputs of the MHW algorithm did not differ significantly on average when as much as 50% missing data were present. It is not however recommended to consider the outputs of time series with this much missing data to be comparable to outputs from an optimal time series. This is because the number of events detected in the time series with high amounts will differ greatly. The overall metrics of the events may be comparable between the time series, but the actual events detected will be different. A simple correction for missing data in a time series is to linearly interpolate over the gaps. It is not however recommended to do this with more than 40% missing data as this begins to dramatically distort the algorithms ability to compute metrics for individual MHWs. If this is necessary to do for some reason, the resultant MHWs for the entire time series can be used to infer the chronic and acute stress that organisms may face in a given location, but any individual events detected should not be taken as an accurate recording.

The decadal trends in times series very rapidly affect the creation of climatologies. That being said, normal ranges of decadal trends (e.g. 0.1 – 0.5C/dec) do not have a significant effect on the detection of MHW metrics. Furthermore, the effect of decadal trends is very predictable and when taken with time series length and the year in which an event in question has occurred it is possible to infer a correction for the maximum intensity. The effect this has on the duration can also be worked out by considering the general raise (or fall) in the mean climatology and how that may engulf neighbouring days or even other events. A concept to consider with the increase in duration from added decadal trends is that the temperatures in the time series increase “faster” than the 90th percentile threshold. So as the decadal trend increases, the MHW effectively spreads outwards. If the rate of onset/decline for the MHW was more gradual (e.g. the NWA event) it will increase in duration more rapidly. If the rate of onset/decline was more rapid (e.g. the Med event), then the duration of the MHW won’t change much with a larger decadal trend. If MHWs have close neighbours then as they spread outward they may encounter and be engulfed into one another. This reduces the overall count of the MHWs detected in a time series while increasing the apparent duration of the events.

Conclusions

We have shown here that researchers must not shy away from the use of sub-optimal time series when the situation calls for it, such as coastal research or sub-surface analyses. Time series length may have an unpredictable effect on MHW results, but this may be corrected for within reason, and we have shown that time series lengths as short as 10 years are still useful for MHW research. Any shorter than 10 years however and the relationship between time series length and the effect on MHW metrics becomes too unpredictable to provide any corrections with confidence. Missing data has a larger effect, but is less of a concern as linear interpolation can largely fix the challenges this creates, up to a threshold of 40% missing data. Lastly, the errors introduced by long-term trends in the data are the most predictable and when taken with time series length may be corrected for as well. The MHW detection algorithm is very robust and we have shown here that one may be confident in the inter-comparability of ones results when using time series within a generous range of sub-optimal data challenges.

References

- Baumgartner, TR. 1992. "Reconstruction of the History of the Pacific Sardine and Northern Anchovy Populations over the Past Two Millenia from Sediments of the Santa Barbara Basin, California." *CalCOFI Rep* 33: 24–40.
- Hobday, Alistair J, Lisa V Alexander, Sarah E Perkins, Dan A Smale, Sandra C Straub, Eric CJ Oliver, Jessica A Benthuysen, et al. 2016. "A Hierarchical Approach to Defining Marine Heatwaves." *Progress in Oceanography* 141. Elsevier: 227–38.
- Hobday, Alistair J, Eric CJ Oliver, Alex Sen Gupta, Jessica A Benthuysen, Michael T Burrows, Markus G Donat, Neil J Holbrook, et al. 2018. "Categorizing and Naming Marine Heatwaves." *Oceanography* 31 (2). JSTOR: 162–73.
- Oliver, Eric CJ, Markus G Donat, Michael T Burrows, Pippa J Moore, Dan A Smale, Lisa V Alexander, Jessica A Benthuysen, et al. 2018. "Longer and More Frequent Marine Heatwaves over the Past Century." *Nature Communications* 9 (1). Nature Publishing Group: 1324.
- Organization, World Meteorological. 2011. *Guide to Climatological Practices*. World Meteorological Organization (WMO).
- . 2017. *WMO Guidelines on the Calculation of Climate Normals*. World Meteorological Organization (WMO).
- Pachauri, Rajendra K, Leo Meyer, J van Ypersele, Sander Brinkman, Line van Kesteren, Noémie Leprince-Ringuet, and Fijke Van Boxmeer. 2014. "Climate Change 2014 Synthesis Report."
- Philander, S George H. 1983. "El Nino Southern Oscillation Phenomena." *Nature* 302 (5906). Nature Publishing Group: 295.
- Reynolds, Richard W, Thomas M Smith, Chunying Liu, Dudley B Chelton, Kenneth S Casey, and Michael G Schlax. 2007. "Daily high-resolution-blended analyses for sea surface temperature." *Journal of Climate* 20 (22). NOAA, Natl Climatic Data Ctr, Asheville, NC 28801 USA: 5473–96.
- Salinger, J., A.J. Hobday, R.J. Matear, T.J. O'Kane, J.S. Risbey, P. Dunstan, J.P. Eveson, et al. 2016. "Chapter One - Decadal-Scale Forecasting of Climate Drivers for Marine Applications." In, edited by Barbara E. Curry, 74:1–68. Advances in Marine Biology. Academic Press. <https://doi.org/https://doi.org/10.1016/bs.amb.2016.04.002>.
- Schlegel, Robert W, and Albertus J Smit. 2018. "HeatwaveR: A Central Algorithm for the Detection of Heatwaves and Cold-Spells." *The Journal of Open Source Software* 3: 821.
- Smale, Dan A, Thomas Wernberg, Eric CJ Oliver, Mads Thomsen, Ben P Harvey, Sandra C Straub, Michael T Burrows, et al. 2019. "Marine Heatwaves Threaten Global Biodiversity and the Provision of Ecosystem Services." *Nature Climate Change*. Nature Publishing Group, 1.
- Wasserstein, Ronald L., Allen L. Schirm, and Nicole A. Lazar. 2019. "Moving to a World Beyond 'P<0.05'." *The American Statistician* 73 (sup1). Taylor & Francis: 1–19. <https://doi.org/10.1080/00031305.2019.1583913>.
- Zhao, Zijie, and Maxime Marin. 2019. "A Matlab Toolbox to Detect and Analyze Marine Heatwaves." *The Journal of Open Source Software* 4: 1124.